

func.new

Ben L. Titzer, Oct 2025 CG Meeting

2025-10-29

Proposal in a nutshell

[intro]

Proposal in a nutshell

[verse]

In the old days

We could just flip a bit

T.L.B. shutdown (OS kernel)

And new instructions were ready

All the baddies came (cash money)

And took over

And that was bad (privilege escalation)

Proposal in a nutshell

[chorus]
func.new (ooh-ooh)
is all we need
func.new (ooh-ooh)
is all we need
new code
with just the caps we want (no escalation)

Proposal in a nutshell

[verse]

So hey we did Wah-zum

Now there's no R...C...Es

Code and data izza sep-ah-rut

But sometimes

Interpreters go slow (oh oh oh)

A whole new module is a hassle

Proposal in a nutshell

```
[chorus]  
func.new (ooh-ooh)  
is all we need  
func.new (ooh-ooh)  
is all we need  
new code  
with just the caps we want (no escalation)
```

Proposal in a nutshell

[verse]

Inside a module, it's all our show

But the tools rearrange us

So better let them know

What the new code uses

Proposal in a nutshell

[chorus]
func.new (ooh-ooh)
is all we need
func.new (ooh-ooh)
is all we need
new code
with just the caps we want (no escalation)

Proposal in a nutshell

[interlude]

Proposal in a nutshell

[chorus]
func.new (ooh-ooh)
is all we need
func.new (ooh-ooh)
is all we need
new code
with just the caps we want (no escalation)

Proposal in a nutshell

[verse]

So we propose
a new mechanism
for new code
but no funny business

Background

- Wasm has separate code and data, aka Harvard architecture
 - Cannot change code at runtime (or even read it)
 - Important security properties:
 - Control-flow integrity (partly due to virtualized execution stack)
 - No remote code execution (RCE)
 - Important analysis properties:
 - Closed-world assumption: can analyze all of a module's access to internals
 - wasm-opt: remove dead code and data from a module

Background

- Wasm has separate code and data, aka Harvard architecture
 - Cannot change code at runtime (or even read it)
 - Important security properties:
 - Control-flow integrity (partly due to virtualized execution stack)
 - No remote code execution (RCE)
 - Important analysis properties:
 - Closed-world assumption: can analyze all of a module's access to internals
 - `wasm-opt`: remove dead code and data from a module
- **New code? That's a host capability!**
 - Wasm MVP launched with JavaScript APIs to make new modules
 - `new WebAssembly.Module(bytes)`
 - `WebAssembly.compile(bytes)`
 - `WebAssembly.compileStreaming(bytes)`

Motivation for the proposal

- Guest virtual machines: a VM in a VM
 - Python, JS on Wasm (yes that's a thing!), Lua, C#, Java, CPU emulators
 - Interpreter performance is bad: **10-100X** slower than JITing
 - Interpreters on Wasm also typically **2-4X** slower than on native
 - Is partial evaluation of an interpreter (Futamura) a solution?


Motivation for the proposal

- Guest virtual machines: a VM in a VM
 - Python, JS on Wasm (yes that's a thing!), Lua, C#, Java, CPU emulators
 - Interpreter performance is bad: **10-100X** slower than JITing
 - Interpreters on Wasm also typically **2-4X** slower than on native
 - Is partial evaluation of an interpreter (Futamura) a solution?
- Existing solutions on the Web
 - JS APIs require entire module at a time
 - Issues with synchronous compile limits

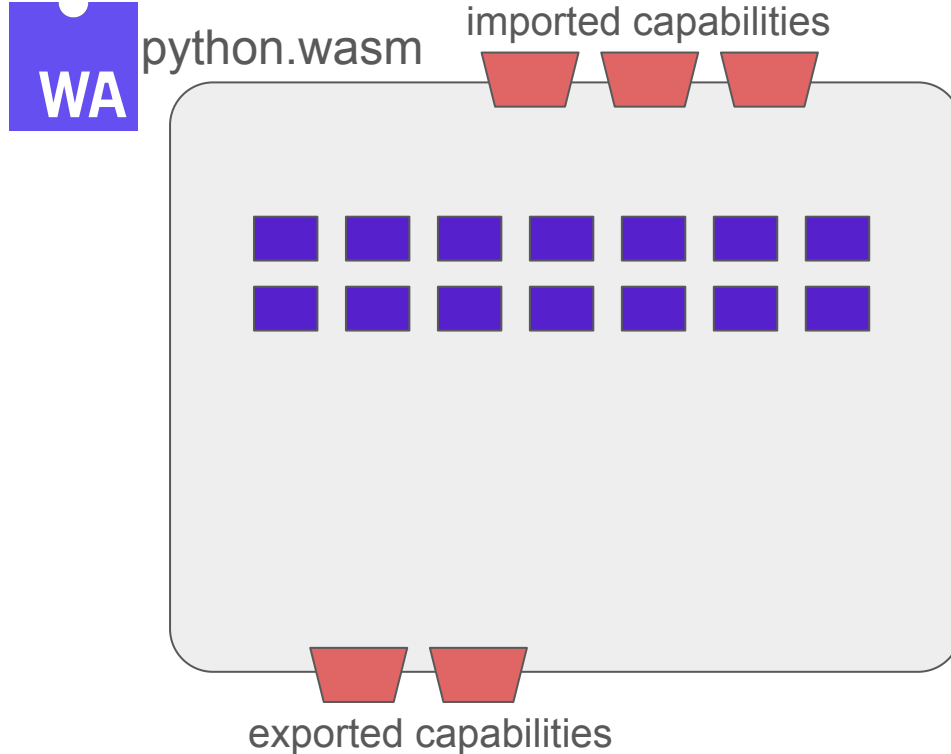
Motivation for the proposal

- Guest virtual machines: a VM in a VM
 - Python, JS on Wasm (yes that's a thing!), Lua, C#, Java, CPU emulators
 - Interpreter performance is bad: **10-100X** slower than JITing
 - Interpreters on Wasm also typically **2-4X** slower than on native
 - Is partial evaluation of an interpreter (Futamura) a solution?
- Existing solutions on the Web
 - JS APIs require entire module at a time
 - Issues with synchronous compile limits
- Existing solutions elsewhere
 - Wasm engine embedding API (`wasm-c-api`)
 - Engine-specific embedding APIs, e.g. WAMR, wasm3, wasmtime
 - Super-powered host function: `wizeng.new_funcref(bytes)`
 - => No portable solution yet

Motivating Example - Guest language runtime

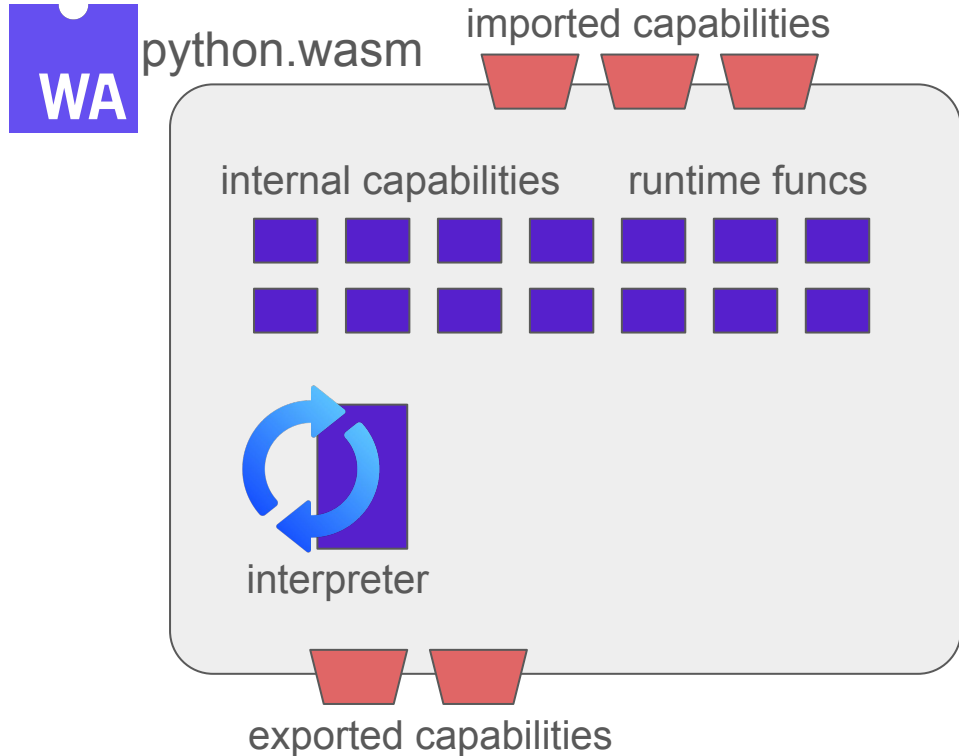


Motivating Example - Guest language runtime



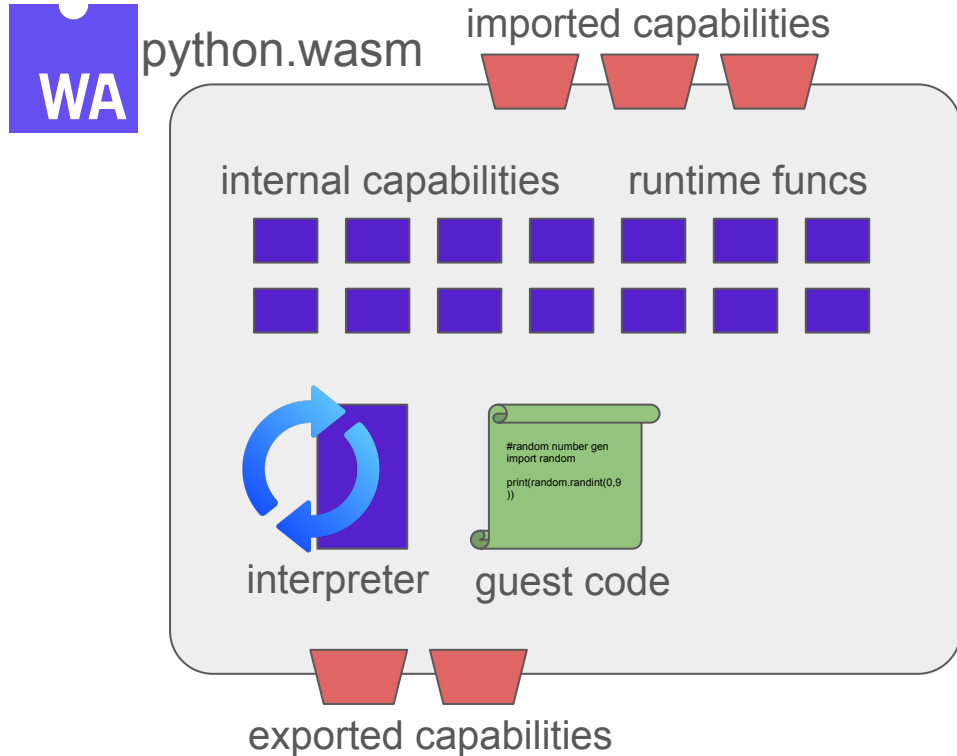
- What's inside?
 - Just Wasm code!

Motivating Example - Guest language runtime



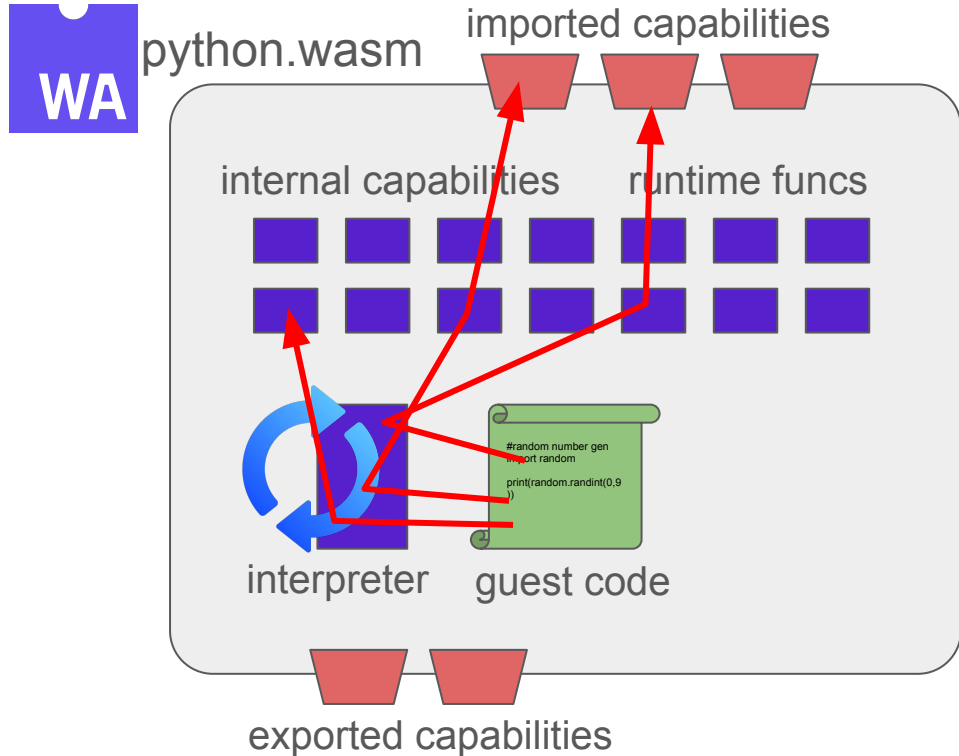
- What's inside?
 - Just Wasm code!
 - A division between internal runtime and an interpreter

Motivating Example - Guest language runtime




- What's inside?
 - Just Wasm code!
 - A division between internal runtime and an interpreter
 - That happens to be Turing-complete

Motivating Example - Guest language runtime



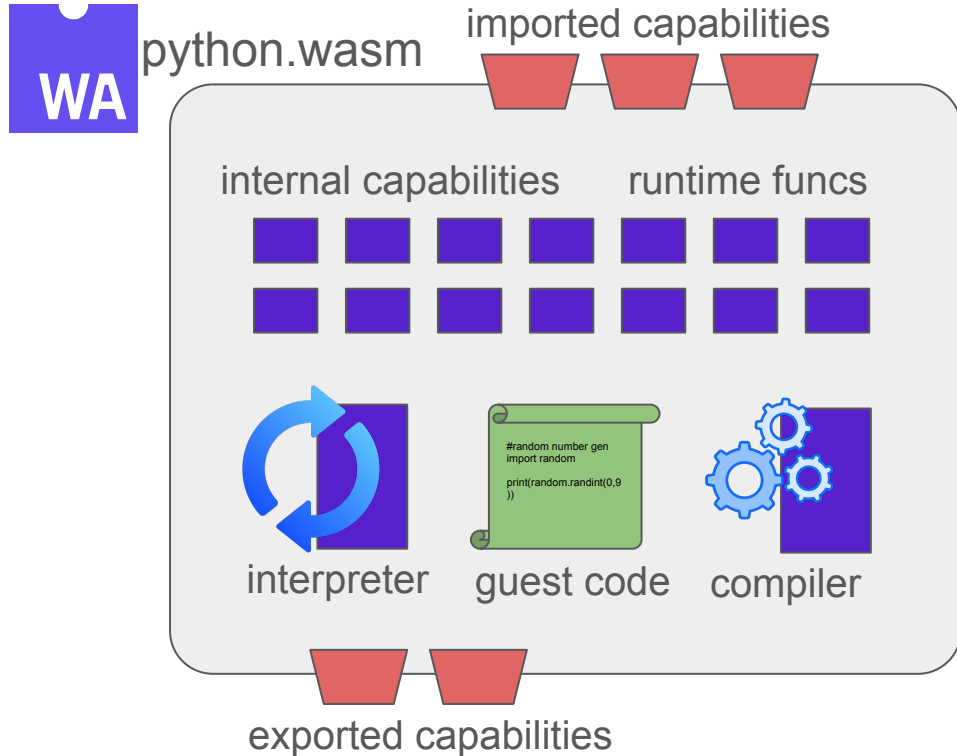
- What's inside?
 - Just Wasm code!
 - A division between internal runtime and an interpreter
 - That happens to be Turing-complete
 - Through the interpreter loop, guest can “drive” the runtime’s internal and imported capabilities.

Motivating Example - Guest language runtime



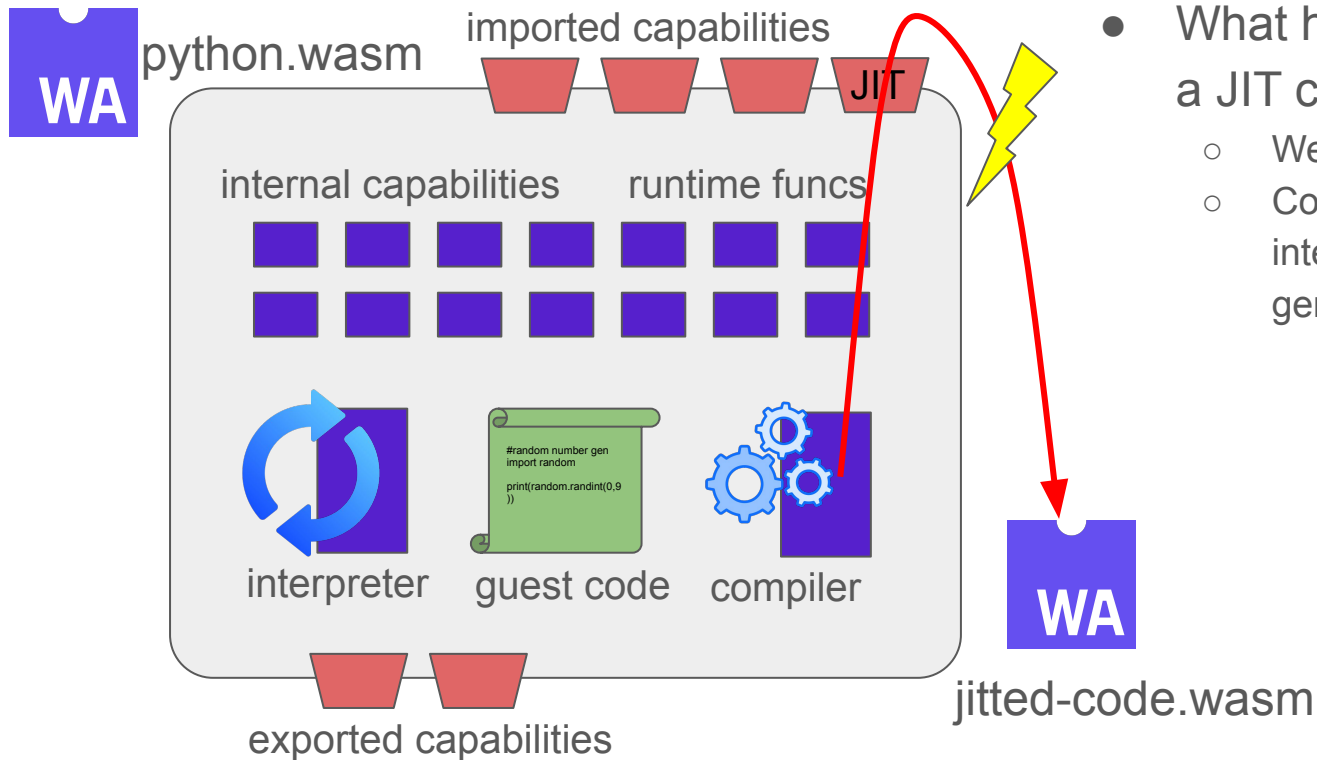
- What's inside?
 - Just Wasm code!
 - A division between internal runtime and an interpreter
 - That happens to be Turing-complete
 - Through the interpreter loop, guest can “drive” the runtime’s internal and imported capabilities.
 - But interpreted code cannot break the encapsulation of its module.

Motivating Example - Guest language runtime



- What happens when we add a JIT compiler?

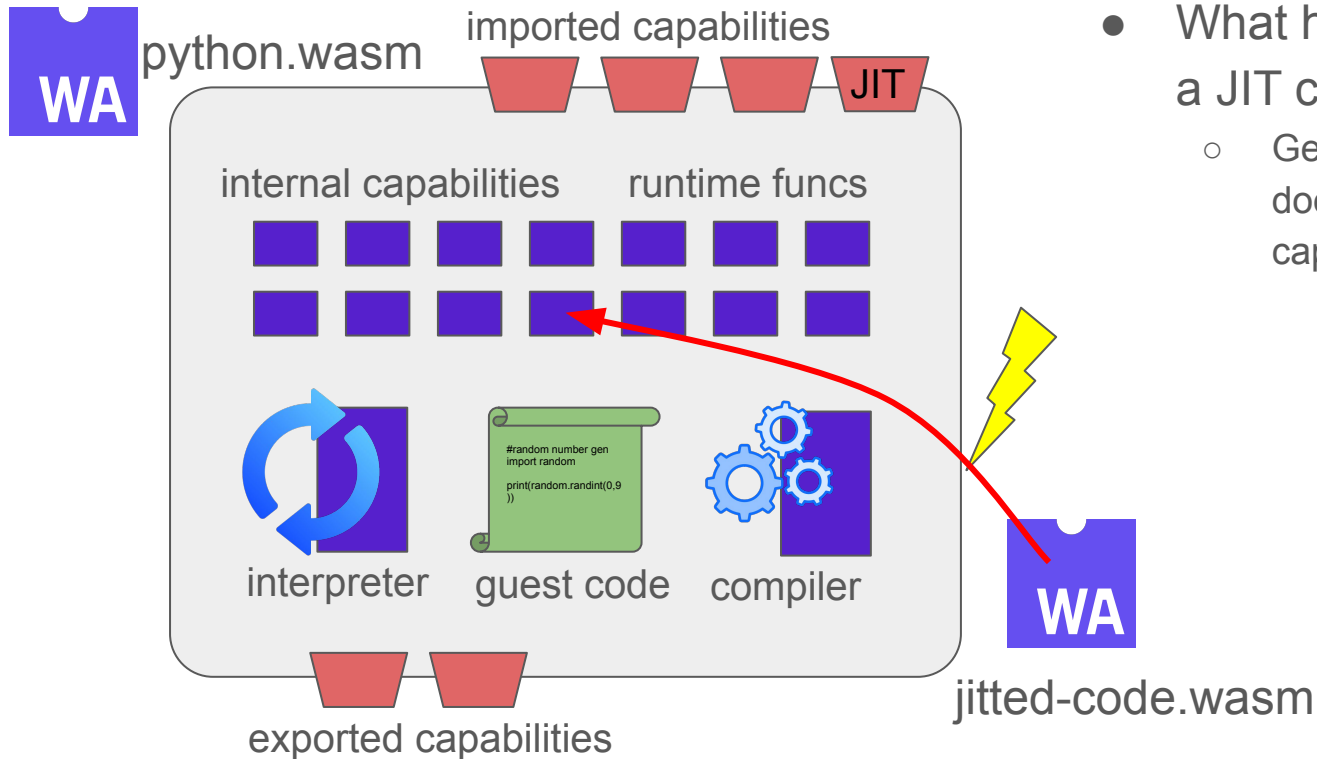
Motivating Example - Guest language runtime



- What happens when we add a JIT compiler?

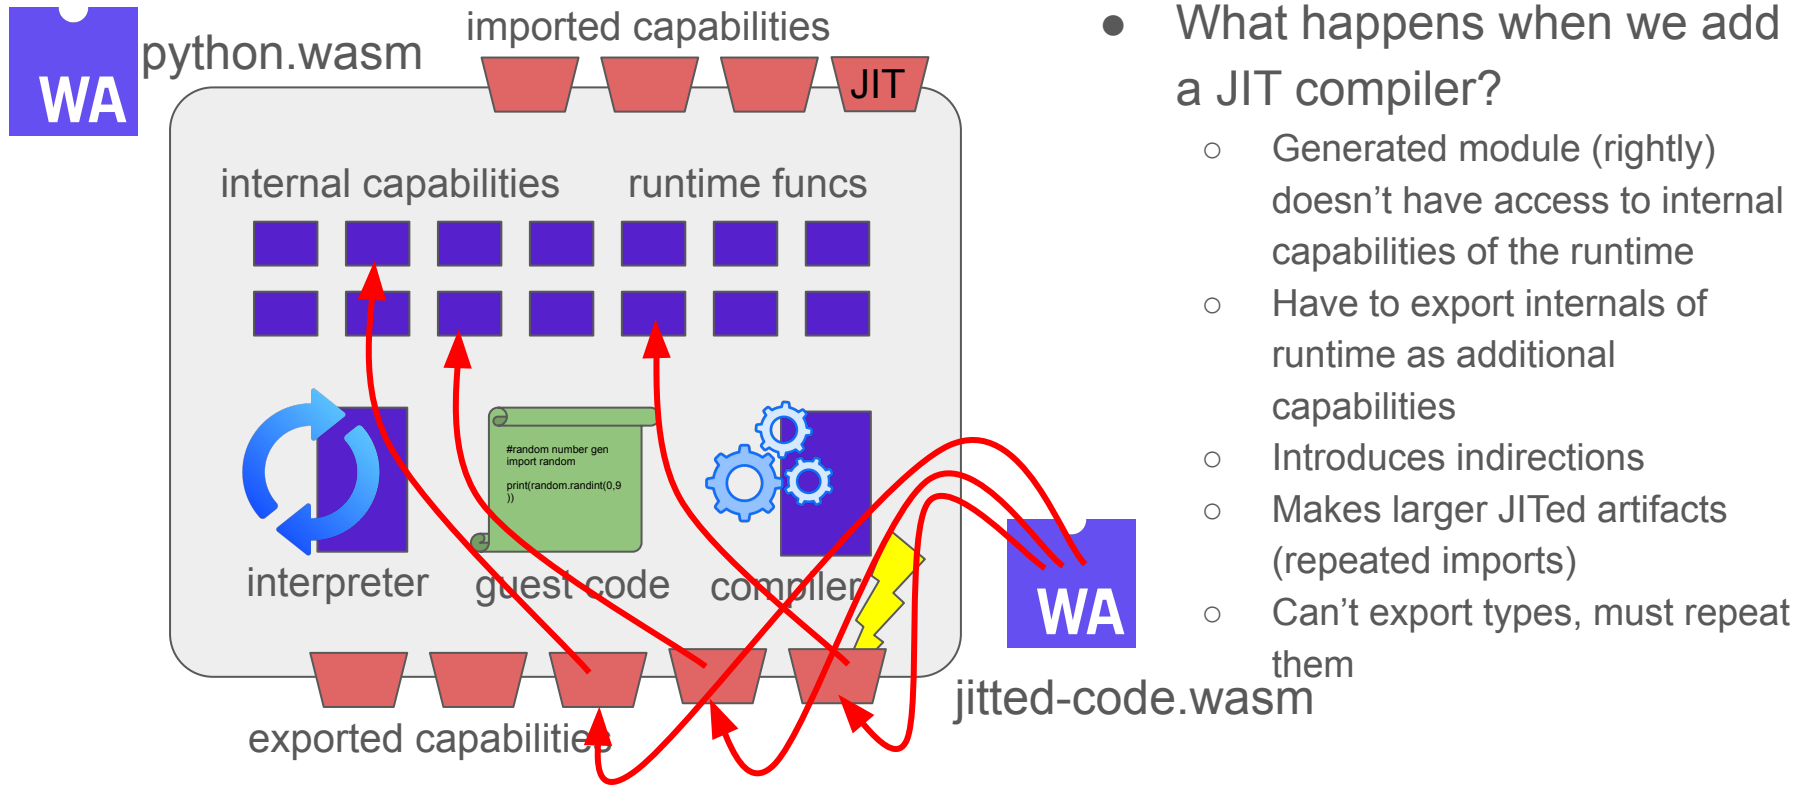
- We need a new host capability
- Could have accidental interposition or corruption on generated code

Motivating Example - Guest language runtime

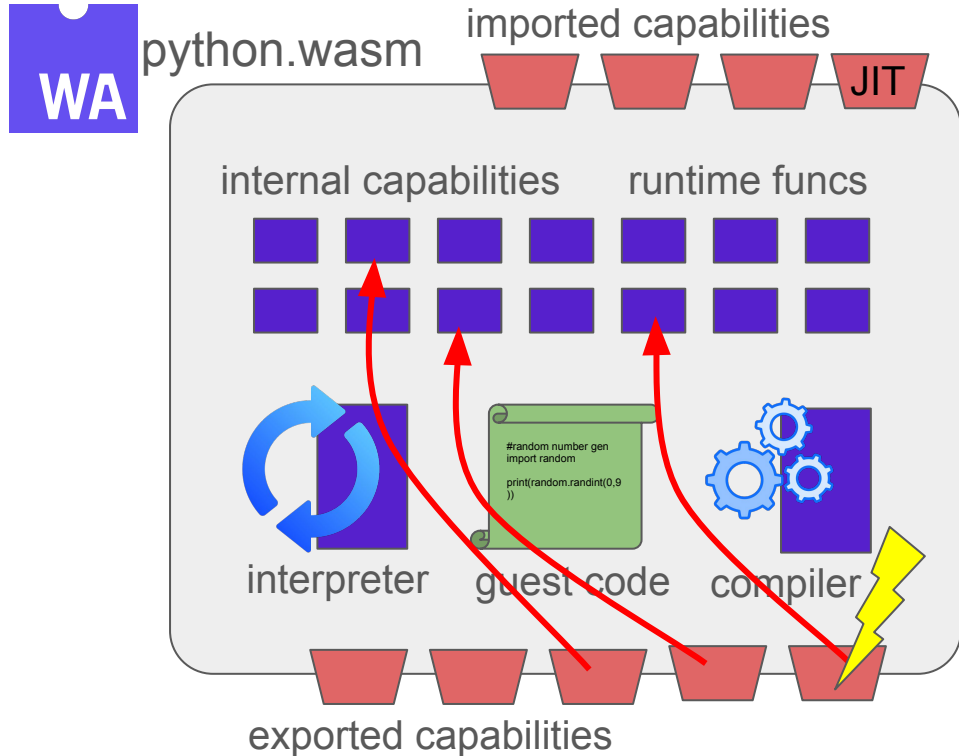


- What happens when we add a JIT compiler?
 - Generated module (rightly) doesn't have access to internal capabilities of the runtime

Motivating Example - Guest language runtime



Motivating Example - Guest language runtime



- What happens when we add a JIT compiler?
 - Generated module (rightly) doesn't have access to internal capabilities of the runtime
 - Have to export internals of runtime as additional capabilities
 - Breaks modularity of the guest language runtime

Proposal: fine-grained JIT interface

- A new bytecode: `func.new $mt $ft $env (start, end)`
 - `$mt`: memory index
 - `$ft`: function type index
 - `$env`: environment index
 - `(start, end)`: indexes into memory where bytecode is stored
- An engine executes `func.new` by:
 - Copying bytes from `(start, end)` in the memory `$mt`
 - Validating the code as if the body of a function with signature `$ft`
 - Returns a new funcref of type `(ref null $ft)` upon success

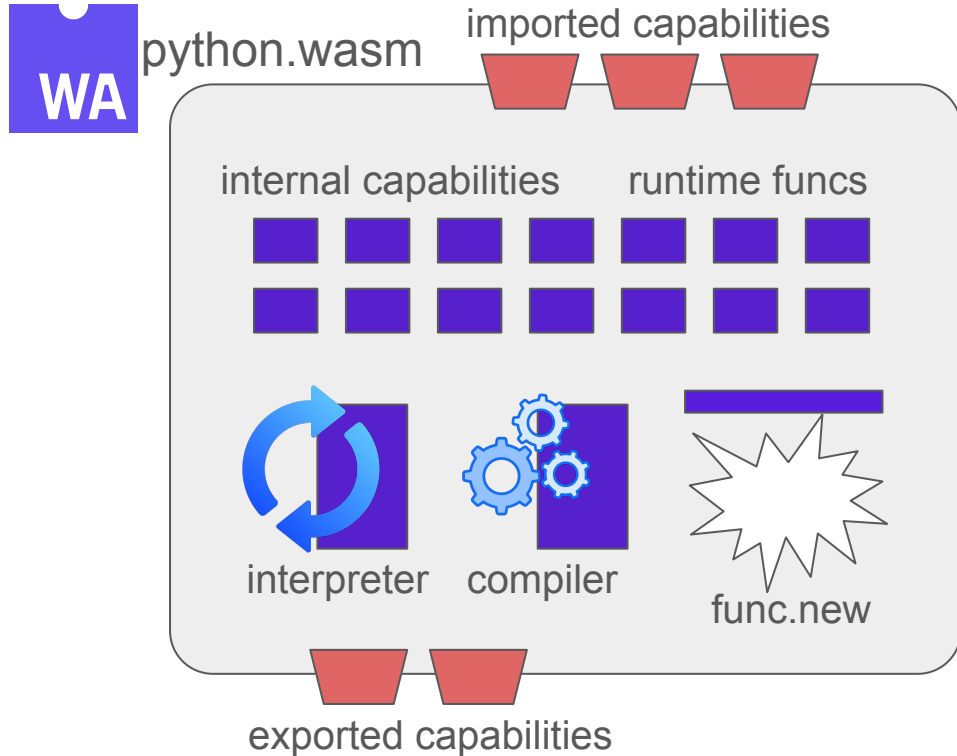
Proposal: fine-grained JIT interface

- A new bytecode: `func.new $mt $ft $env (start, end)`
 - `$mt`: memory index
 - `$ft`: function type index
 - `$env`: environment index
 - `(start, end)`: indexes into memory where bytecode is stored
- A new section: *environments* (scopes)
 - Declares a subset of the encapsulating module to be accessible to new code
 - Renumbers declarations starting from 0

Proposal: fine-grained JIT interface

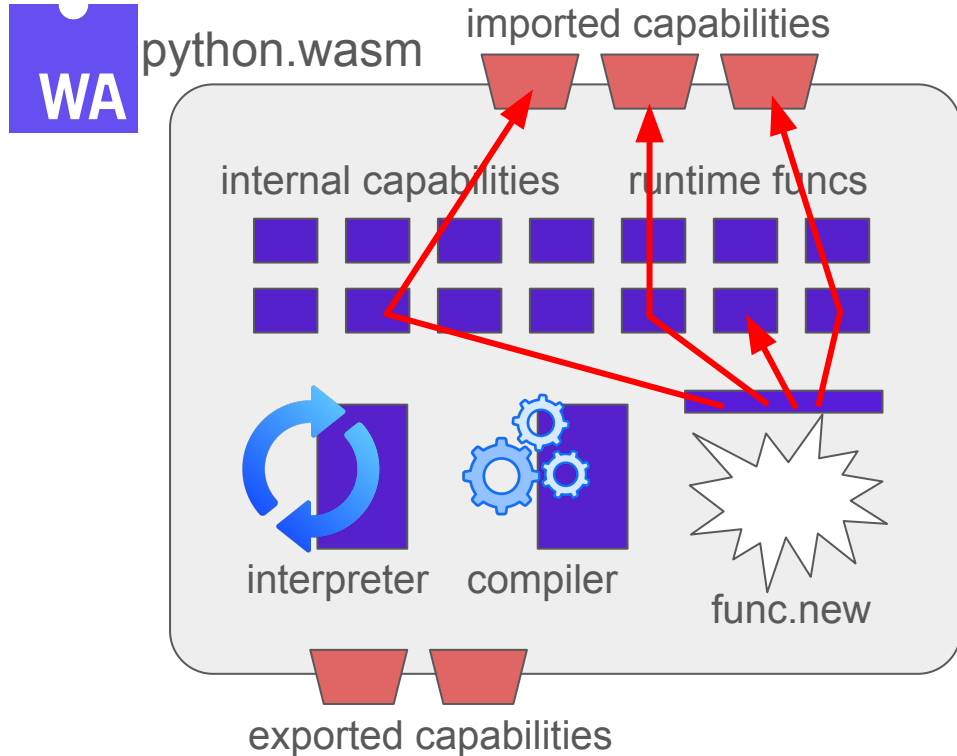
- A new bytecode: `func.new $mt $ft $env (start, end)`
 - `$mt`: memory index
 - `$ft`: function type index
 - `$env`: environment index
 - `(start, end)`: indexes into memory where bytecode is stored
- A new section: *environments* (scopes)
 - Declares a subset of the encapsulating module to be accessible to new code
 - Renumbers declarations starting from 0
- A flag for memories: `code`
 - Similar to `shared` flag, allows a memory to be used for `func.new`
 - Prevents accidental use of a memory to make code

Motivating Example - Guest language runtime



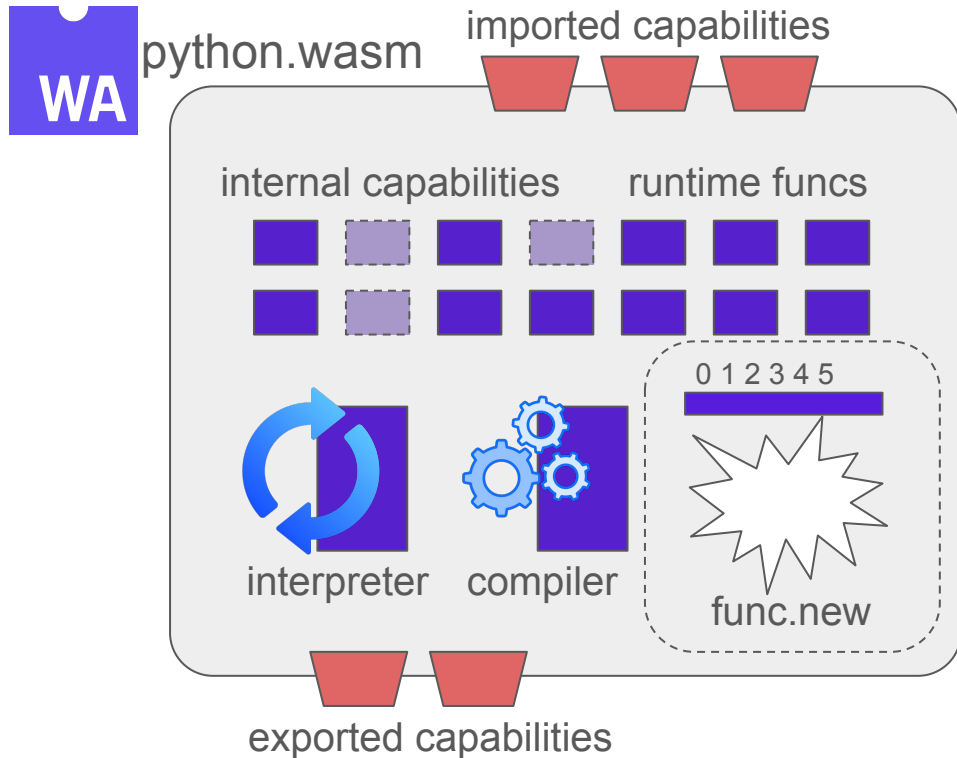
- func.new plus its environment is a controlled “hole” in the module for dynamically creating new functions

Motivating Example - Guest language runtime



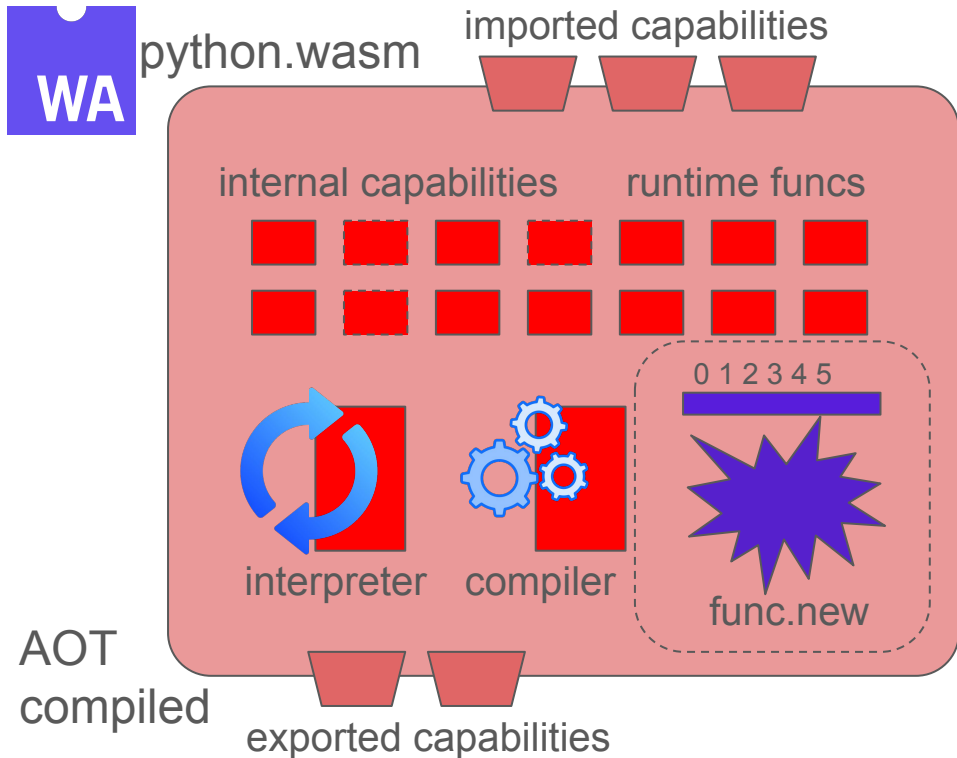
- What does `func.new` accomplish?
 - No new host capability needed for new code
 - New code has no additional privileges
 - Preserves guest runtime module encapsulation
 - Environment describes exactly what new code can use from the containing module, allowing static reasoning
 - Engine has simplified code validator (not whole module)

Motivating Example - Guest language runtime



- What does `func.new` accomplish?
 - Toolchains can still reorganize modules without requiring new code to renumber
 - Allows sound DCE, inlining, module combining, and other transformations

Motivating Example - Guest language runtime



- What does func.new accomplish?
 - Requires engine to support parsing/validating **only function bodies**, not whole modules

Dynamically
interpreter / compiled

Example .wat Usage

```
(module
  (type $t1 (func)) ;; the type for new functions
  (func $f1 ...)
  (func $f2 ...)
  (func $f3 ...)
  (memory $m1 code 1 1) ;; the memory used to temporarily store code for func.new
  (memory $m2 1 1)      ;; a memory accessible to new code

  (env| $s1              ;; the scope a new function may use
    (func $f1 $f2)       ;; expose $f1 and $f2 to new code
    (memory $m2))        ;; expose only $m2 to new code

  (func $gen
    (local $n (ref $t1)) ;; a variable to hold the new funcref
    ...
    (local.set $n
      (func.new $m1 $t1 $s1 ;; code lives in $m1, result sig is $t1, scope is $s1
        (i32.const 1024) (i32.const 10))) ;; code is stored at address 1024 and is 10 bytes long
    ...
    (call_ref $t1 (local.get $n)) ;; call the new function!!
  )
)
```

Example .wat Usage

Code flag for memory

```
(module
  (type $t1 (func)) ;; the type for new functions
  (func $f1 ...)
  (func $f2 ...)
  (func $f3 ...)
  (memory $m1 code 1 1) ;; the memory used to temporarily store code for func.new
  (memory $m2 1 1) ;; a memory accessible to new code

  (env| $s1 ;; the scope a new function may use
    (func $f1 $f2) ;; expose $f1 and $f2 to new code
    (memory $m2)) ;; expose only $m2 to new code

  (func $gen
    (local $n (ref $t1)) ;; a variable to hold the new funcref
    ...
    (local.set $n
      (func.new $m1 $t1 $s1 ;; code lives in $m1, result sig is $t1, scope is $s1
        (i32.const 1024) (i32.const 10))) ;; code is stored at address 1024 and is 10 bytes long
    ...
    (call_ref $t1 (local.get $n)) ;; call the new function!!
  )
)
```

Example .wat Usage

Environment for new code

```
(module
  (type $t1 (func)) ;; the type for new functions
  (func $f1 ...)
  (func $f2 ...)
  (func $f3 ...)
  (memory $m1 code 1 1) ;; the memory used to temporarily store code for func.new
  (memory $m2 1 1)      ;; a memory accessible to new code

  (env $s1                ;; the scope a new function may use
    (func $f1 $f2)        ;; expose $f1 and $f2 to new code
    (memory $m2))         ;; expose only $m2 to new code

  (func $gen
    (local $n (ref $t1)) ;; a variable to hold the new funcref
    ...
    (local.set $n
      (func.new $m1 $t1 $s1                ;; code lives in $m1, result sig is $t1, scope is $s1
        (i32.const 1024) (i32.const 10))) ;; code is stored at address 1024 and is 10 bytes long
    ...
    (call_ref $t1 (local.get $n))          ;; call the new function!!
  )
)
```

Example .wat Usage

func.new bytecode usage

```
(module
  (type $t1 (func)) ;; the type for new functions
  (func $f1 ...)
  (func $f2 ...)
  (func $f3 ...)
  (memory $m1 code 1 1) ;; the memory used to temporarily store code for func.new
  (memory $m2 1 1)      ;; a memory accessible to new code

  (env| $s1              ;; the scope a new function may use
    (func $f1 $f2)       ;; expose $f1 and $f2 to new code
    (memory $m2))        ;; expose only $m2 to new code

  (func $gen
    (local $n (ref $t1)) ;; a variable to hold the new funcref
    ...
    (local.set $n
      (func.new $m1 $t1 $s1 ;; code lives in $m1, result sig is $t1, scope is $s1
        (i32.const 1024) (i32.const 10))) ;; code is stored at address 1024 and is 10 bytes long
    ...
    (call_ref $t1 (local.get $n)) ;; call the new function!!
  )
)
```

Discussion points

- Is this nested modules? (component model)
- Asynchronous compilation
 - Different opcode `func.new_async`?
 - Blocking by default?
 - Non-blocking, blocks on call?
 - `Func.ready` allows polling?
- How to handle failure, e.g. resource exhaustion
- Feature testing
- How to incorporate custom sections (debug names, branch/compilation hints, other annotations)
- Is the code flag for memories useful?
- Should we accept Wasm GC array of bytes?
- Multiple functions at a time? (for mutual direct calls)
- What about new GC types? (`type.new`)

Poll for Phase 1

Entry requirements

- There is general interest within the CG in this feature.
- The CG believes the feature is in-scope and will plausibly be workable.

Func.new is All We Need

<https://suno.com/song/19e0679a-81a2-419f-8cc7-ff85ac49c23d>

[verse]

In the old days
We could just flip a bit
T.L.B. shutdown (OS kernel)
And new instructions were ready
All the baddies came (cash money)
And took over
And that was bad (privilege escalation)

[chorus]

func.new (ooh-ooh)
is all we need
func.new (ooh-ooh)
is all we need
new code
with just the caps we want (no escalation)

[verse]

So hey we did Wah-zum
Now there's no R...C...Es
Code and data izza sep-ah-rut
But sometimes
Interpreters go slow (oh oh oh)
A whole new module is a hassle

[verse]

Inside a module, it's all our show
But the tools rearrange us
So better let them know
What the new code uses

[verse]

So we propose
a new mechanism
for new code
but no funny business