



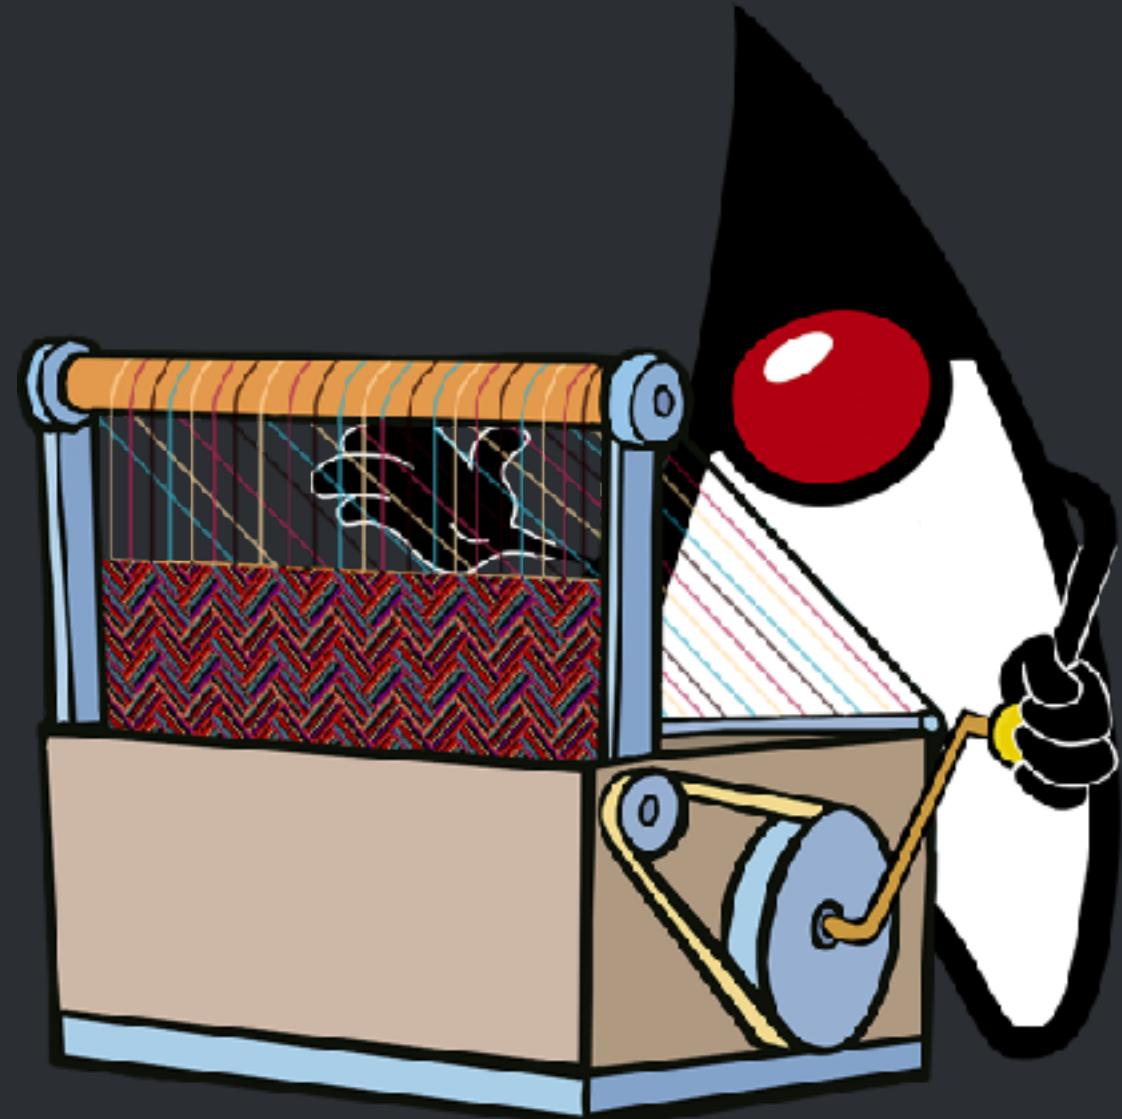
# The Design of User-Mode Threads In Java

---

Ron Pressler

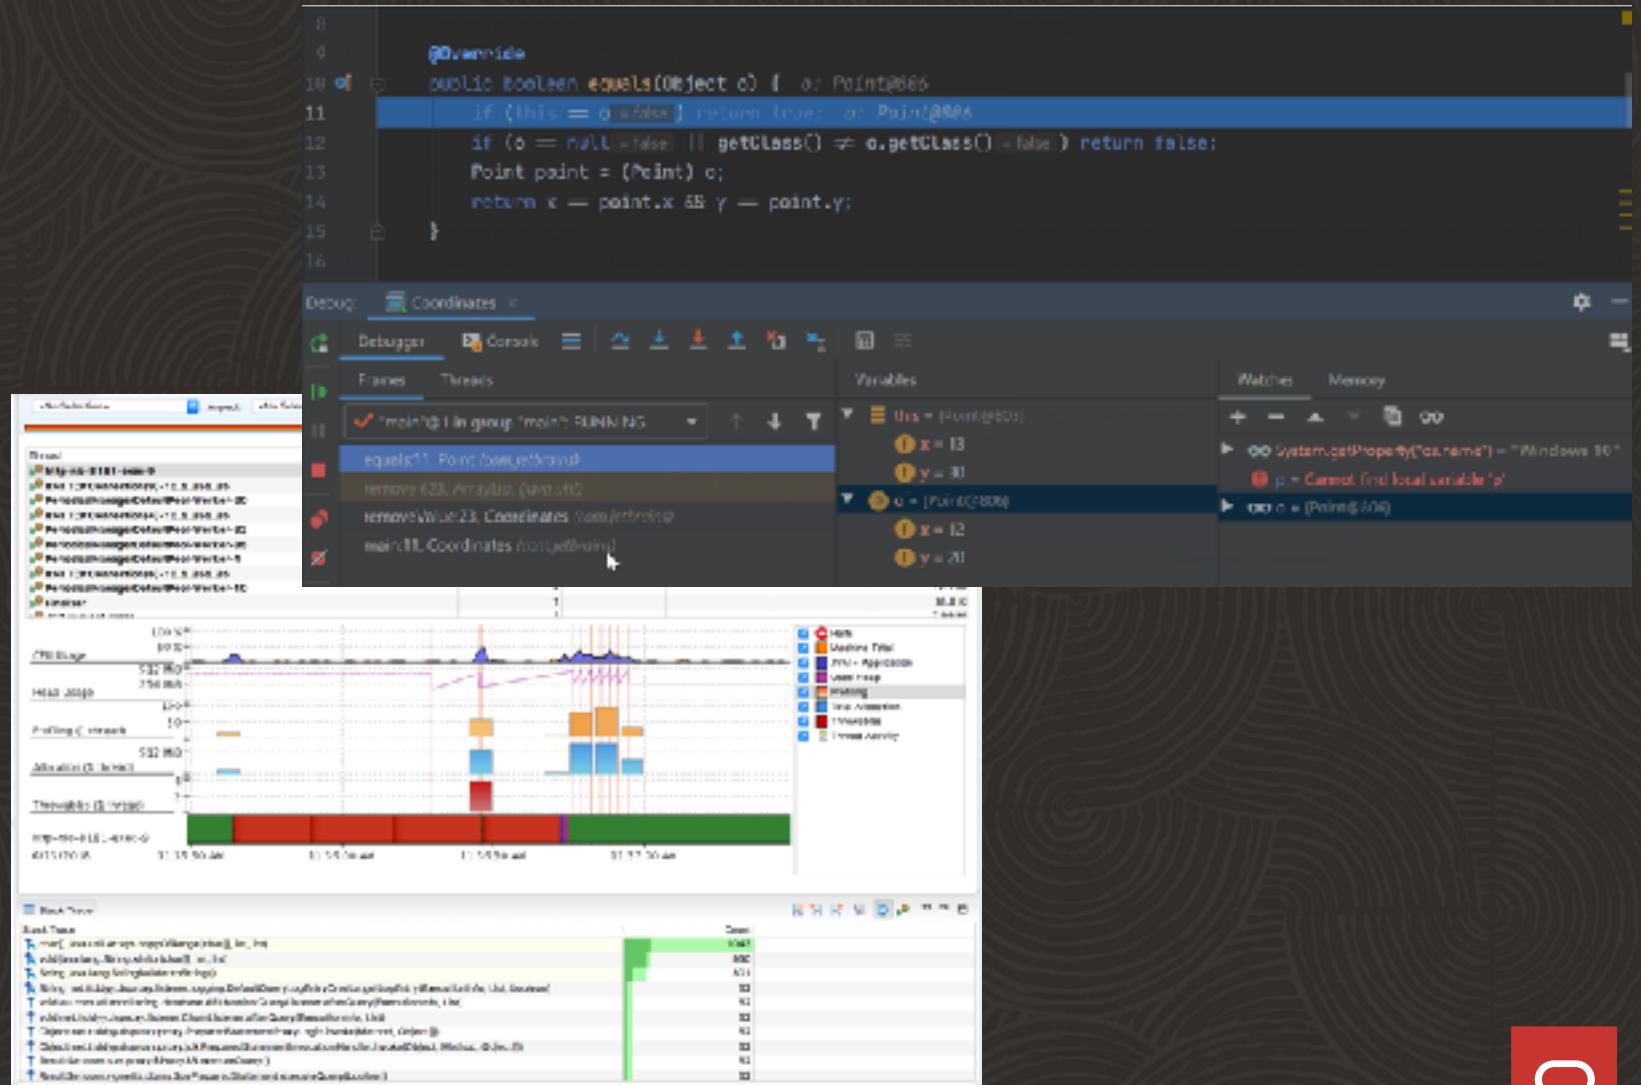
Java Platform Group

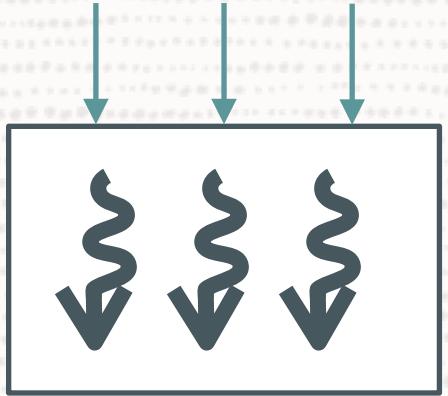
January 25, 2021



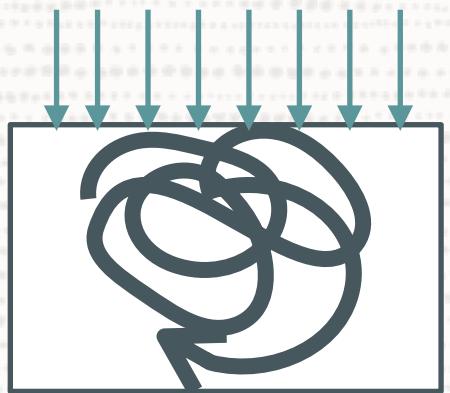
# Java Is Made of Threads

- Exceptions
  - Thread Locals
  - Debugger
  - Profiler (JFR)





simple  
less scalable



scalable,  
complex,  
non-interoperable,  
hard to debug/profile

SYNC

Programmer  
OS / Hardware



OR

ASYNC

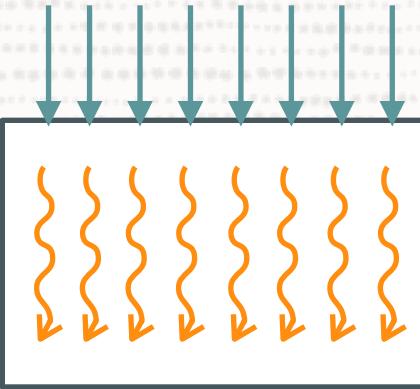
Programmer  
OS / Hardware



*Codes Like sync, Scales Like Async*

Connections

App



Programmer

OS / Hardware



“We must carefully balance  
**conservation** and **innovation**”

— Mark Reinhold

- **Forward Compatibility:** we want existing code to enjoy new functionality
- We want to **correct past mistakes** and **start afresh**

“The solutions of **yesterday**  
are the problems of **today**”

— Brian Goetz



# Threads *in Java*

- The use of `Thread.currentThread()` and `ThreadLocal` is pervasive. Without support, or with changed behaviour, little existing code would run.
- Other parts are superseded by new APIs since Java 5 so their datedness/clunkiness is mostly hidden/ignored.

# *Threads in Java*

- `java.lang.Thread`
- The Java runtime is well positioned to implement threads.
- Resizable stacks (possible b/c we only need to support Java).
- Context-switching in user-mode.
- Pluggable schedulers, default optimised for transaction processing.

# *Threads in Java*

When code in a virtual thread calls an I/O method in the JDK,  
suspend the virtual thread,  
start a non-blocking I/O operation in the OS,  
the scheduler schedules another virtual thread,  
when I/O completes re-submit waiting thread to scheduler.

Module java.base  
Package java.util.concurrent

## Class ConcurrentHashMap<K,V>

java.lang.Object  
  java.util.AbstractMap<K,V>  
    java.util.concurrent.ConcurrentHashMap<K,V>

### Type Parameters:

K - the type of keys maintained by this map

V - the type of mapped values

### All Implemented Interfaces:

Serializable, ConcurrentMap<K,V>, Map<K,V>

```
public class ConcurrentHashMap<K, V>
extends AbstractMap<K, V>
implements ConcurrentMap<K, V>, Serializable
```

A hash table supporting full concurrency of retrievals and high expected concurrency for updates. This class obeys the same functional specification as `Hashtable`, and includes versions of methods corresponding to each method of `Hashtable`. However, even though all operations are thread-safe, retrieval operations do not entail locking, and there is *not* any support for locking the entire table in a way that prevents all access. This class is fully interoperable with `Hashtable` in programs that rely on its thread safety but not on its synchronization.

Module java.base  
Package java.nio.channels

## Class SocketChannel

java.lang.Object
  java.nio.channels.spi.AbstractInterruptibleChannel
    java.nio.channels.SelectableChannel
      java.nio.channels.spi.AbstractSelectableChannel
        java.nio.channels.SocketChannel

### All Implemented Interfaces:

Closeable, AutoCloseable, ByteChannel, Channel, GatheringByteChannel, InterruptibleChannel, NetworkChannel, ReadableByteChannel, ScatteringByteChannel, WritableByteChannel

```
public abstract class SocketChannel
extends AbstractSelectableChannel
implements ByteChannel, ScatteringByteChannel, GatheringByteChannel, NetworkChannel
```

A selectable channel for stream-oriented connecting sockets.

A socket channel is created by invoking one of the `open` methods of this class. It is not possible to create a channel for an arbitrary, pre-existing socket. A newly-created socket channel is open but not yet connected. An attempt to invoke an I/O operation upon an unconnected channel will cause a `NotYetConnectedException` to be thrown. A socket channel can be connected by invoking its `connect` method; once connected, a socket channel remains connected until it is closed. Whether or not a socket channel is connected may be determined by invoking its

Module java.base  
Package java.util.concurrent.locks

## Class ReentrantLock

java.lang.Object
  java.util.concurrent.locks.ReentrantLock

### All Implemented Interfaces:

Serializable, Lock

```
public class ReentrantLock
extends Object
implements Lock, Serializable
```

A reentrant mutual exclusion lock with the same basic behavior and semantics as the implicit monitor lock accessed using `synchronized` methods and statements, but with extended capabilities.

A `ReentrantLock` is owned by the thread last successfully locking, but not yet unlocking it. A thread invoking `lock` will return, successfully acquiring the lock, when the lock is not owned by another thread. The method will return immediately if the current thread already owns the lock. This can be checked using methods `isHeldByCurrentThread()`, and `getHoldCount()`.

The constructor for this class accepts an optional `fairness` parameter. When set `true`, under contention, locks favor granting access to the longest-waiting thread. Otherwise this lock does not guarantee any particular access order. Programs using fair locks accessed by many threads may display lower overall throughput (i.e., are slower).

Module java.base  
Package java.io

## Class InputStream

java.lang.Object
  java.io.InputStream

### All Implemented Interfaces:

Closeable, AutoCloseable

### Direct Known Subclasses:

AudioInputStream, ByteArrayInputStream, FileInputStream, FilterInputStream, ObjectInputStream, PipedInputStream, SequenceInputStream, StringBufferInputStream

```
public abstract class InputStream
extends Object
implements Closeable
```

This abstract class is the superclass of all classes representing an input stream of bytes.

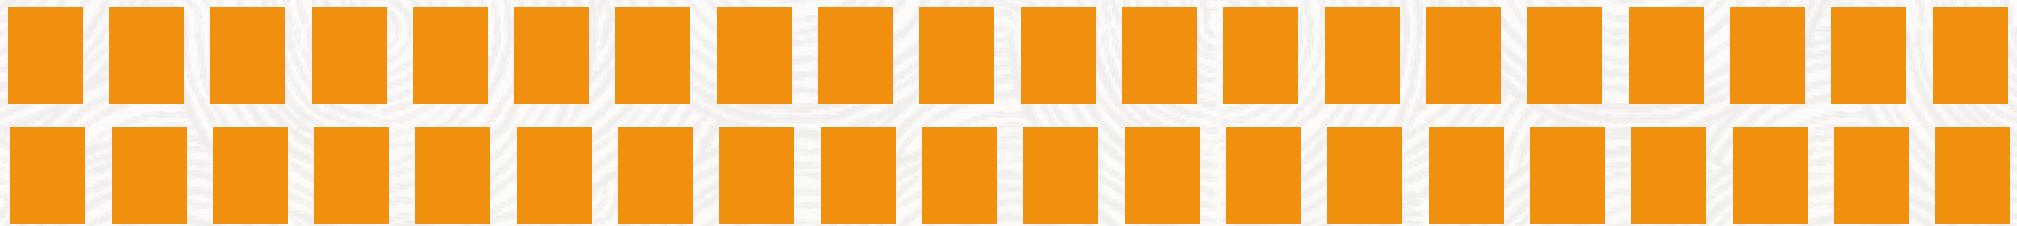
Applications that need to define a subclass of `InputStream` must always provide a method that returns the next byte of input.

### Since:

1.0



virtual threads



“carrier” platform threads managed by a scheduler

# async/await

c#

JavaScript

Kotlin

C++/Rust

# User-Mode Threads

Erlang

Go

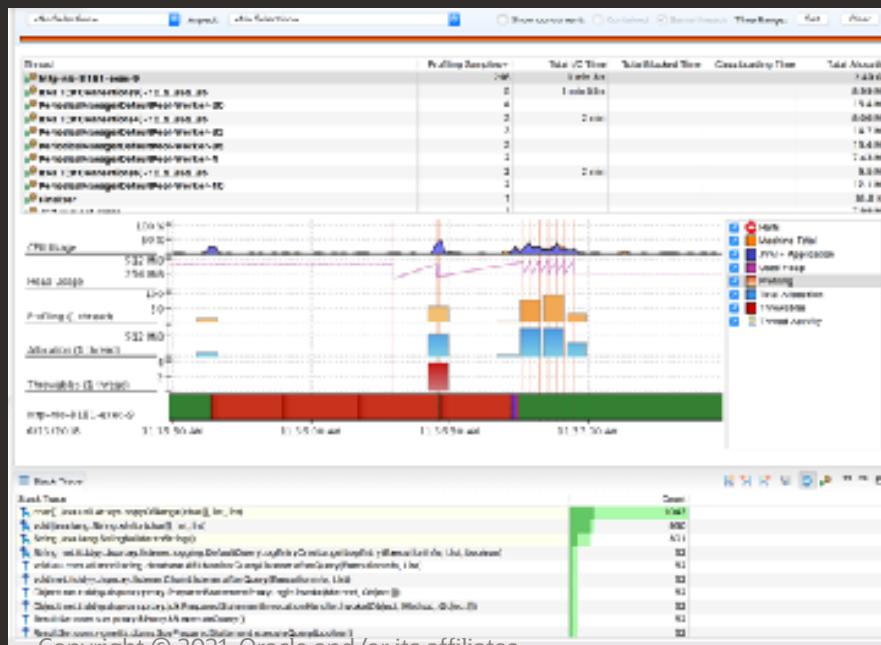
Java

# Process: Unit of Concurrency

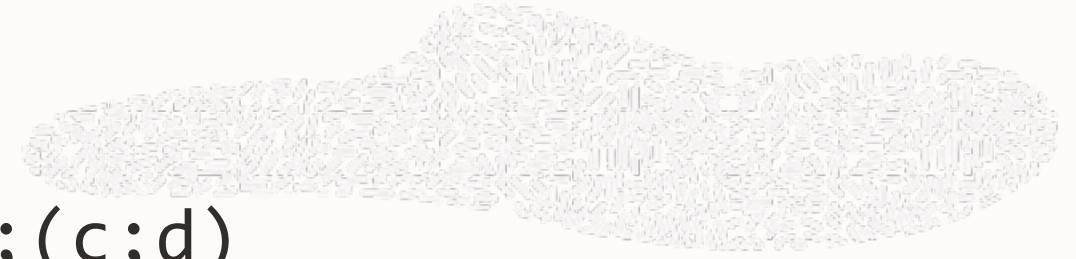
E.g. a transaction



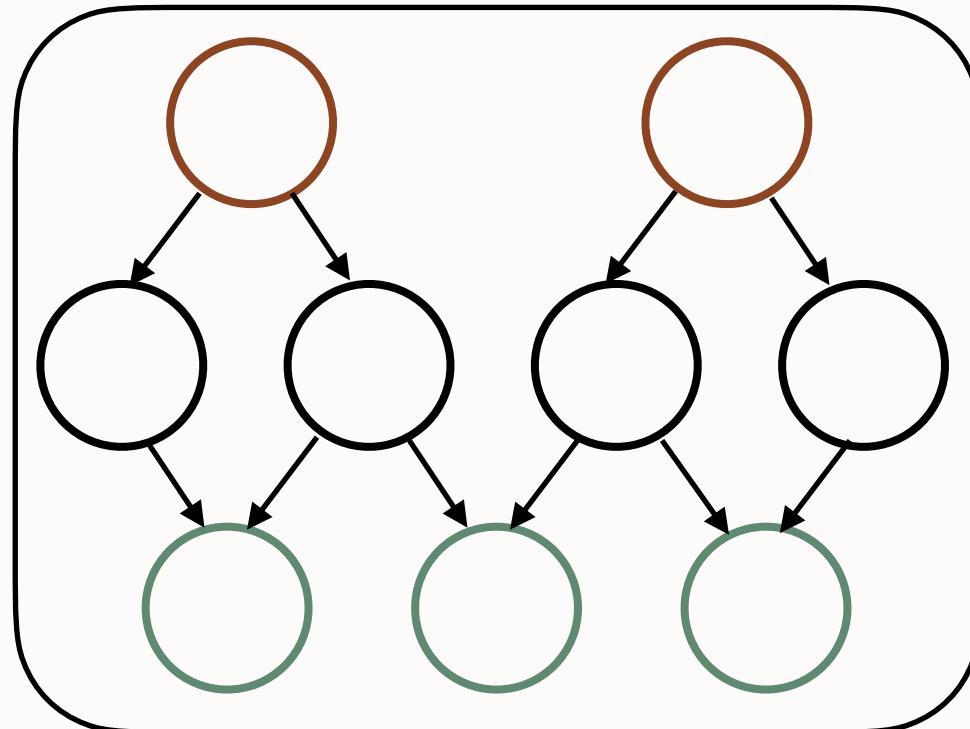
- Code (writing/reading)
  - Troubleshooting: stack traces, debugger single-stepping
  - Profiling



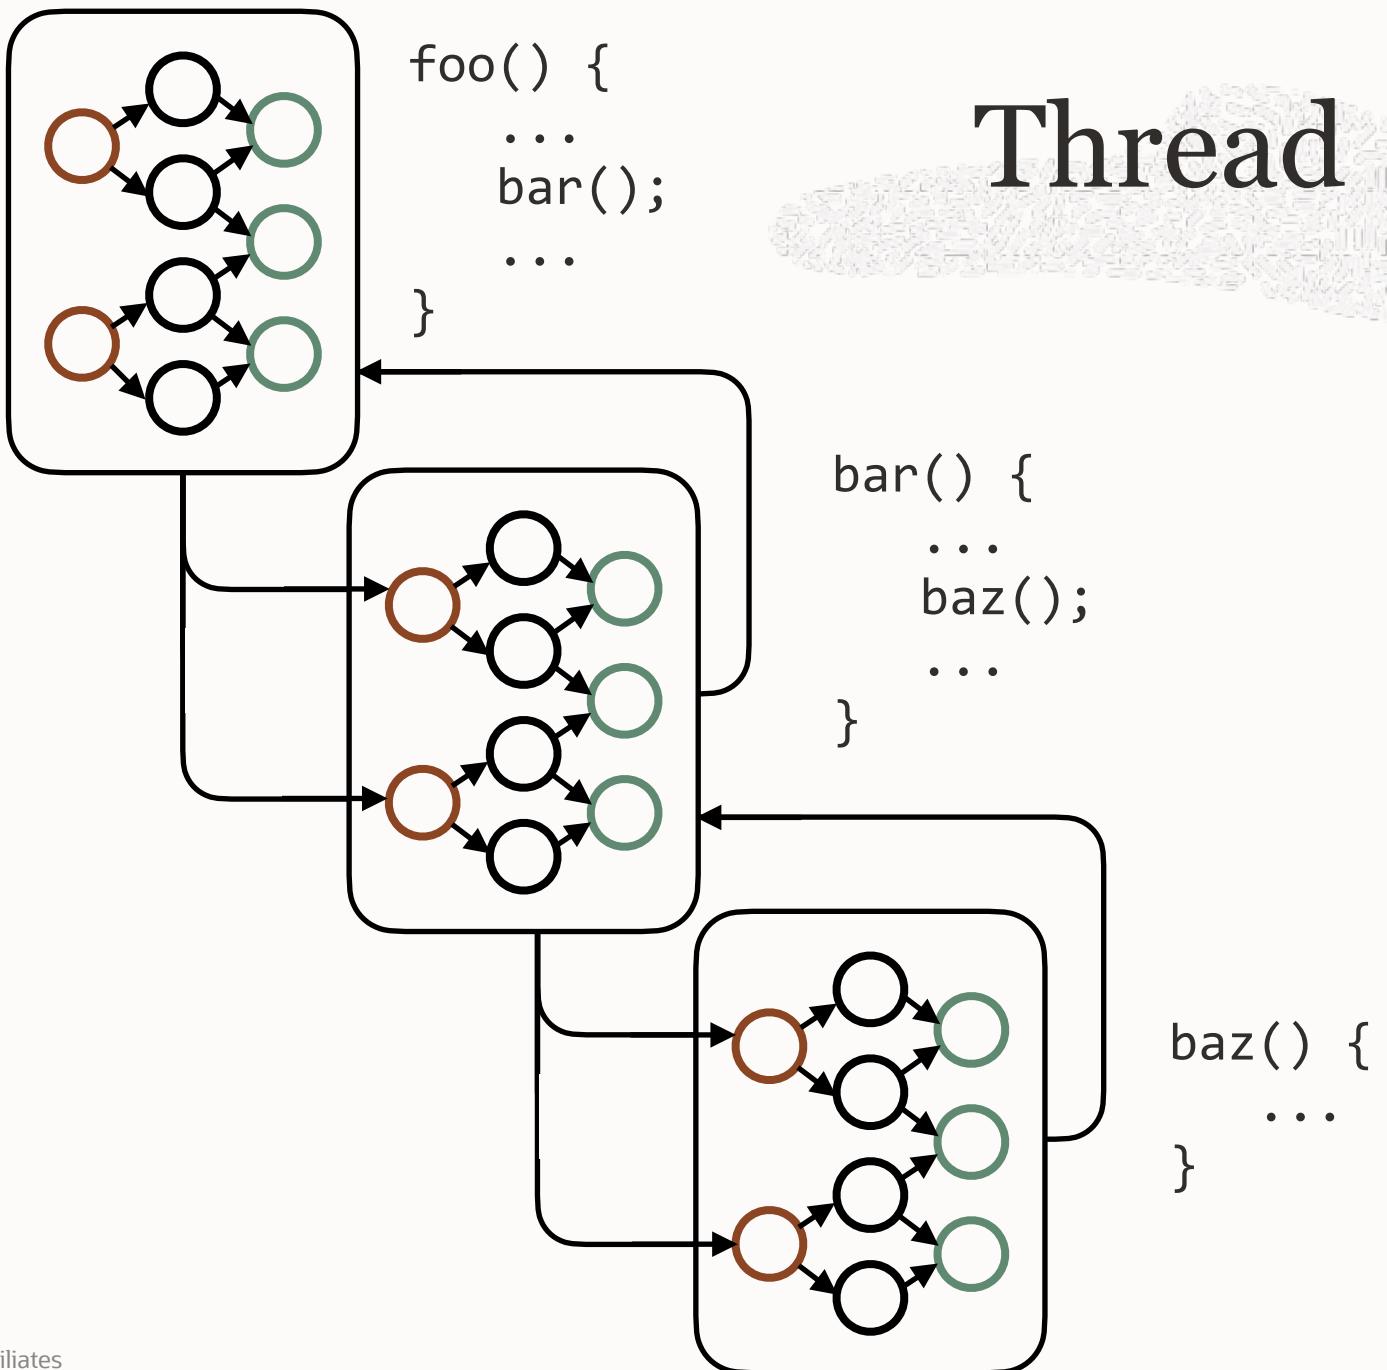
# Process



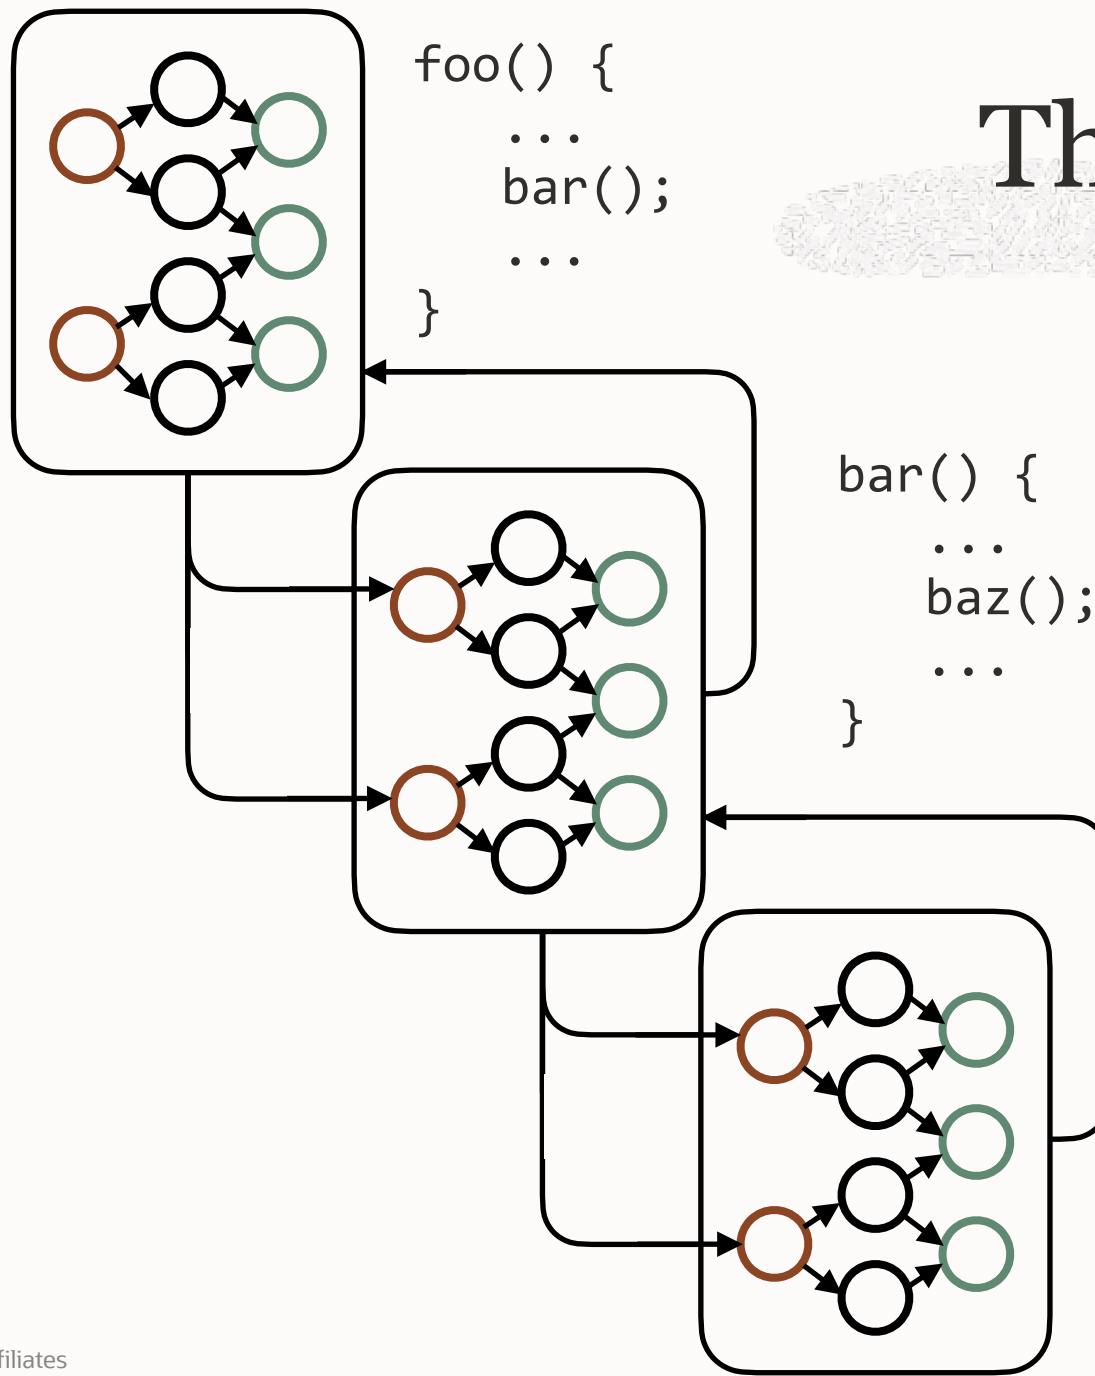
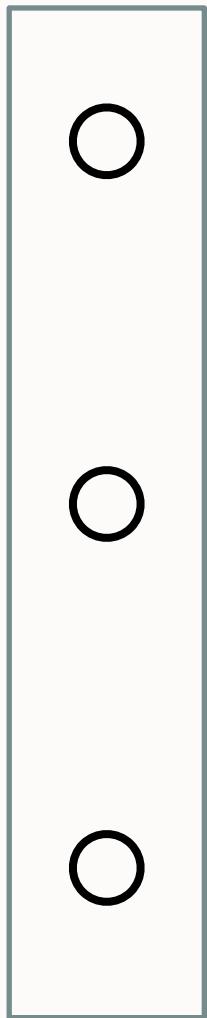
$a;b;c;d = (a;b);(c;d)$



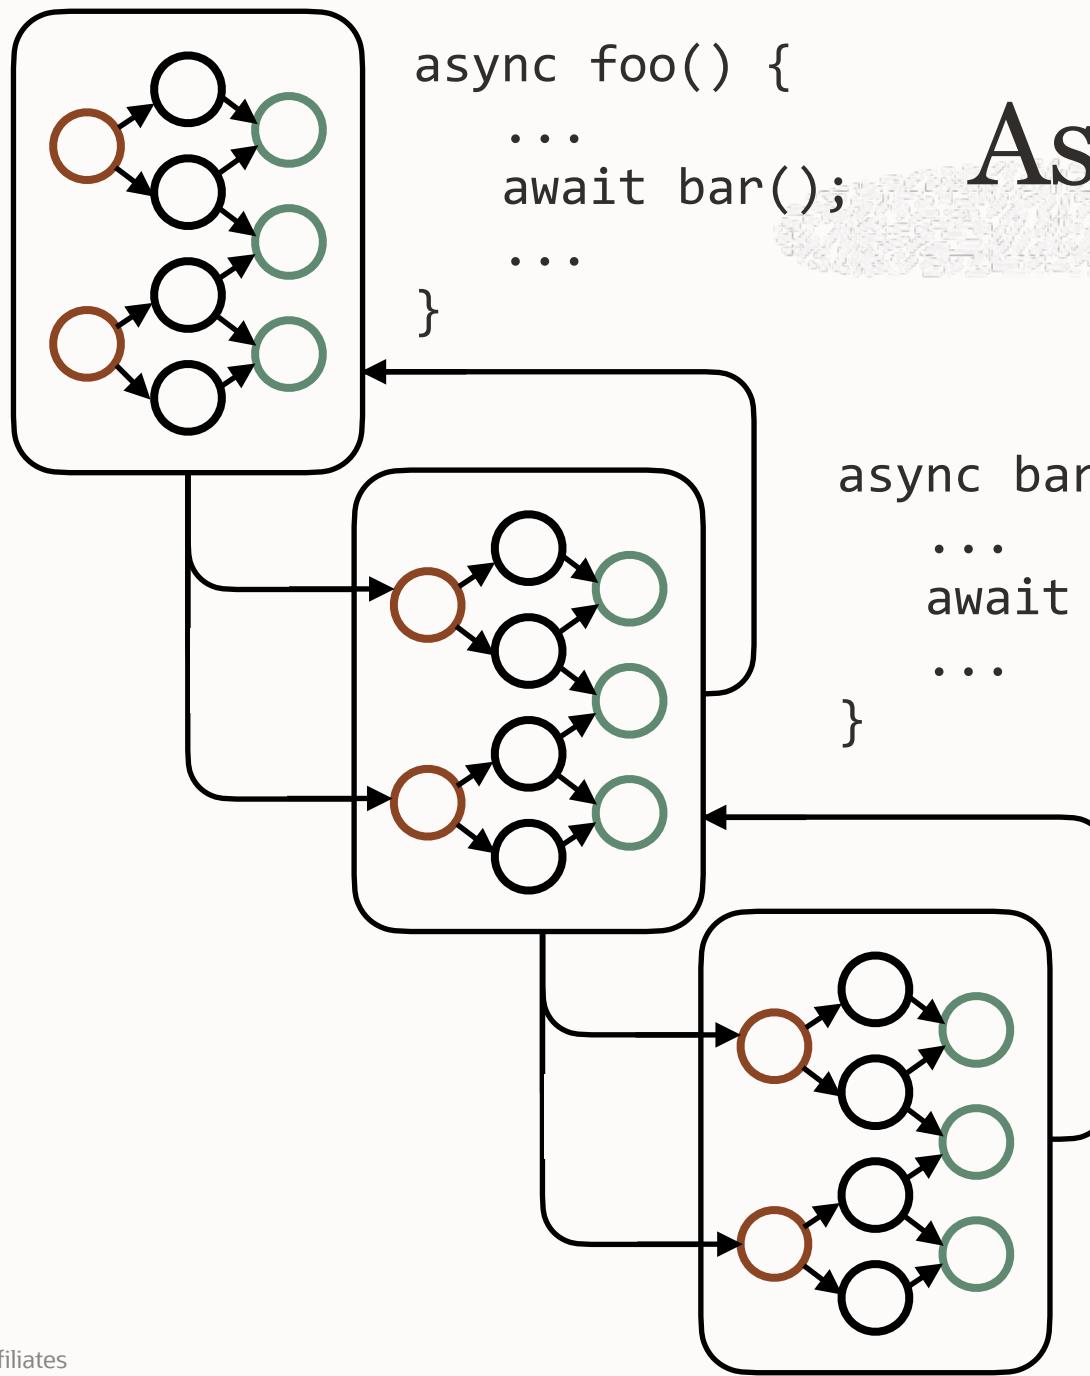
(Nondeterminism <https://youtu.be/9vupFNsND6o> )



Call Stack



# Async/Await



# Thread vs. Async/Await

## Scheduling/interleaving points

Thread: Everywhere *except* where explicitly forbidden (with a CS)

async/await: Nowhere *except* where explicitly allowed (with await)

# Thread vs. Async/Await

## Scheduling/interleaving points

Thread: Everywhere *except* where explicitly forbidden (with a CS)

async/await: Nowhere *except* where explicitly allowed (with await)

# JavaScript

# Thread vs. Async/Await

## Implementation

Thread: Requires integrating with the implementation of subroutines (control over backend)

async/await: Can be implemented in the compiler frontend

# Thread vs. Async/Await

## Implementation

Thread: Requires integrating with the implementation of subroutines (control over backend)

async/await: Can be implemented in the compiler frontend

Kotlin

# Thread vs. Async/Await

## Recursion & virtual calls

Thread: Yes (requires ~~large~~/resizable stacks)

async/await: Can be excluded

# Thread vs. Async/Await

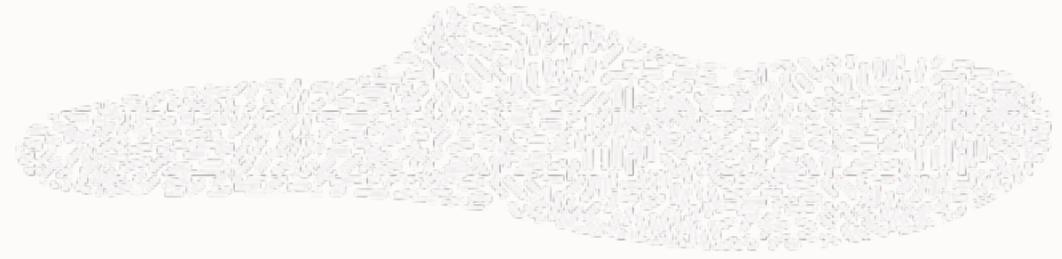
## Recursion & virtual calls

Thread: Yes (requires ~~large~~/resizable stacks)

async/await: Can be excluded

C++/Rust

# Resizable Stack



- Transparent allocation
- Efficient allocation
- No internal pointers/tracked pointers (no FFI)

# Performance

**Latency** — How long an operation takes (s)

**Throughput** — How many operations complete per time unit (ops/s)

**Impact** — How much a metric would improve with full optimisation (%)

# Syntactic Concurrency: Generators et al.



- Updating simulation entities in a frame
- Generators (two processes with an unbuffered channel)

```
def rev_str(my_str):  
    length = len(my_str)  
    for i in range(length - 1, -1, -1):  
        yield my_str[i]  
  
for char in rev_str("hello"):  
    print(char)
```

# Context-Switching Impact: Generators

- Impact: 100%
- Best case latency: ~0ns (monomorphic, fits in cache)

# Concurrency: Transactions

$L = \lambda W$

<https://inside.java/2020/08/07/loom-performance/>

Throughput:  $\lambda = L/W$

Context-switch impact on throughput:  $t/\mu$

$t$  — Mean context-switch latency

$\mu$  — Mean wait (I/O) latency

# Context-Switching Impact: Transactions

- Impact: low if blocking for external events
- Best case latency: 60ns (polymorphic, doesn't fit in cache) (1.5% impact)
- Target latency for  $\leq 5\%$  impact:  $\leq 200\text{ns}$

# Conclusion

- Control over backend
- Rare I/O in FFI
- No internal pointers/pointers tracked
- Efficient and transparent stack resizing
- Threads already in the platform, libraries and tooling
- `async/await` can be implemented on top by languages that want it

# async/await

c#

JavaScript

Kotlin

C++/Rust

# User-Mode Threads

Erlang  
Go  
Java

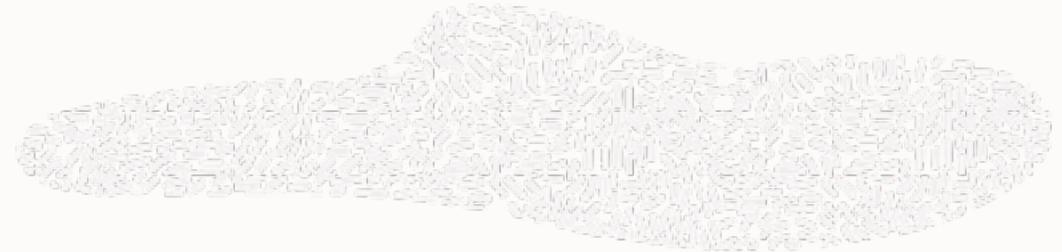
Zig

# Continuations

# Continuation

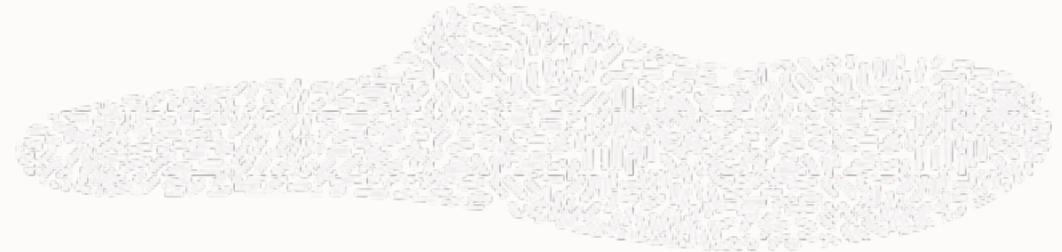
- Delimited
- Stackful
- Asymmetric
- Type **void** (but other types can be trivially implemented on top)
- Scoped (aka multi-prompt, aka “effect handlers”)
- One-shot/non-reentrant (future: cloneable & serializable)

# Continuations



```
class Continuation implements Runnable {  
    public Continuation(ContinuationScope scope, Runnable body);  
    public void run();  
    public boolean isDone();  
    public static void yield(ContinuationScope scope);  
}
```

# Continuations



```
Continuation cont = new Continuation(SCOPE, () -> {
    System.out.println("before");
    Continuation.yield(SCOPE);
    System.out.println("after");
});
```

```
cont.run(); // prints before
cont.isDone(); // false
cont.run(); // prints after
cont.isDone(); // true
```

thread =  
+ continuation  
scheduler

# Thread = Continuation + Scheduler

```
package java.util.concurrent.locks;  
public class LockSupport {  
  
    public static void park(...) {  
        var t = Thread.currentThread();  
        if (t.isVirtual())  
            Continuation.yield(FIBER_SCOPE);  
        else  
            Unsafe.park(false, 0L);  
    }  
  
    public static void unpark(Thread thread) {  
        if (thread.isVirtual()) {  
            t.scheduler.submit(t.continuation);  
        } else {  
            Unsafe.unpark(thread)  
        }  
    }  
}
```

```
var readLock = new java.util.concurrent.locks.ReentrantLock();  
  
readLock.lock();
```

```
java.base/java.lang.Continuation.yield(Continuation.java:396)  
java.base/java.lang.Fiber.maybePark(Fiber.java:597)  
java.base/java.lang.Fiber.park(Fiber.java:525)  
java.base/java.lang.System$2.parkFiber(System.java:2333)  
java.base/jdk.internal.misc.Strands.parkFiber(Strands.java:89)  
java.base/java.util.concurrent.locks.LockSupport.park(LockSupport.java:235)  
java.base/java.util.concurrent.locks.AbstractQueuedSynchronizer.parkAndCheckInterrupt(AbstractQueuedSynchronizer.java:887)  
java.base/java.util.concurrent.locks.AbstractQueuedSynchronizer.acquireQueued(AbstractQueuedSynchronizer.java:919)  
java.base/java.util.concurrent.locks.AbstractQueuedSynchronizer.acquire(AbstractQueuedSynchronizer.java:1242)  
java.base/java.util.concurrent.locks.ReentrantLock.lock(ReentrantLock.java:269)  
Main.lambda$main$0(Main.java:72)  
java.base/java.lang.Fiber.lambda$new$0(Fiber.java:198)  
java.base/java.lang.Continuation.enter(Continuation.java:372)  
java.base/java.lang.Continuation.run(Continuation.java:328)  
java.base/java.lang.Fiber.runContinuation(Fiber.java:366)  
java.base/java.util.concurrent.ForkJoinTask$RunnableExecuteAction.exec(ForkJoinTask.java:1425)  
java.base/java.util.concurrent.ForkJoinTask.doExec(ForkJoinTask.java:290)  
java.base/java.util.concurrent.ForkJoinPool$WorkQueue topLevelExec(ForkJoinPool.java:1017)  
java.base/java.util.concurrent.ForkJoinPool.scan(ForkJoinPool.java:1666)  
java.base/java.util.concurrent.ForkJoinPool.runWorker(ForkJoinPool.java:1599)  
java.base/java.util.concurrent.ForkJoinWorkerThread.run(ForkJoinWorkerThread.java:189)
```

```
var readLock = new java.util.concurrent.locks.ReentrantLock();  
  
readLock.lock();
```



```
java.base/java.lang.Continuation.run(Continuation.java:328)  
java.base/java.lang.Fiber.runContinuation(Fiber.java:366)  
java.base/java.util.concurrent.ForkJoinTask$RunnableExecuteAction.exec(ForkJoinTask.java:1425)  
java.base/java.util.concurrent.ForkJoinTask.doExec(ForkJoinTask.java:290)  
java.base/java.util.concurrent.ForkJoinPool$WorkQueue topLevelExec(ForkJoinPool.java:1017)  
java.base/java.util.concurrent.ForkJoinPool.scan(ForkJoinPool.java:1666)  
java.base/java.util.concurrent.ForkJoinPool.runWorker(ForkJoinPool.java:1599)  
java.base/java.util.concurrent.ForkJoinWorkerThread.run(ForkJoinWorkerThread.java:189)
```

# Thread = Continuation + Scheduler

```
package java.util.concurrent.locks;  
public class LockSupport {  
  
    public static void park(...) {  
        var t = Thread.currentThread();  
        if (t.isVirtual())  
            Continuation.yield(FIBER_SCOPE);  
        else  
            Unsafe.park(false, 0L);  
    }  
  
}  
}
```

```
    public static void unpark(Thread thread) {  
        if (thread.isVirtual()) {  
            t.scheduler.submit(t.continuation);  
        } else {  
            Unsafe.unpark(thread)  
        }  
    }  
}
```

```
var readLock = new java.util.concurrent.locks.ReentrantLock();
readLock.lock();

...
```

```
java.base/java.lang.Continuation.yield(Continuation.java:396)
java.base/java.lang.Fiber.maybePark(Fiber.java:597)
java.base/java.lang.Fiber.park(Fiber.java:525)
java.base/java.lang.System$2.parkFiber(System.java:2333)
java.base/jdk.internal.misc.Strands.parkFiber(Strands.java:89)
java.base/java.util.concurrent.locks.LockSupport.park(LockSupport.java:235)
java.base/java.util.concurrent.locks.AbstractQueuedSynchronizer.parkAndCheckInterrupt(AbstractQueuedSynchronizer.java:887)
java.base/java.util.concurrent.locks.AbstractQueuedSynchronizer.acquireQueued(AbstractQueuedSynchronizer.java:919)
java.base/java.util.concurrent.locks.AbstractQueuedSynchronizer.acquire(AbstractQueuedSynchronizer.java:1242)
java.base/java.util.concurrent.locks.ReentrantLock.lock(ReentrantLock.java:269)
Main.lambda$main$0(Main.java:72)
java.base/java.lang.Fiber.lambda$new$0(Fiber.java:198)
java.base/java.lang.Continuation.enter(Continuation.java:372)
java.base/java.lang.Continuation.run(Continuation.java:328)
java.base/java.lang.Fiber.runContinuation(Fiber.java:366)
java.base/java.util.concurrent.ForkJoinTask$RunnableExecuteAction.exec(ForkJoinTask.java:1425)
java.base/java.util.concurrent.ForkJoinTask.doExec(ForkJoinTask.java:290)
java.base/java.util.concurrent.ForkJoinPool$WorkQueue topLevelExec(ForkJoinPool.java:1017)
java.base/java.util.concurrent.ForkJoinPool.scan(ForkJoinPool.java:1666)
java.base/java.util.concurrent.ForkJoinPool.runWorker(ForkJoinPool.java:1599)
java.base/java.util.concurrent.ForkJoinWorkerThread.run(ForkJoinWorkerThread.java:189)
```

```
var listener = new java.net.ServerSocket(...);  
listener.accept();
```



```
java.base/java.lang.Continuation.yield(Continuation.java:396)  
java.base/java.lang.Fiber.maybePark(Fiber.java:597)  
java.base/java.lang.Fiber.park(Fiber.java:525)  
java.base/java.lang.System$2.parkFiber(System.java:2333)  
java.base/jdk.internal.misc.Strands.parkFiber(Strands.java:89)  
java.base/sun.nio.ch.NioSocketImpl.park(NioSocketImpl.java:183)  
java.base/sun.nio.ch.NioSocketImpl.park(NioSocketImpl.java:212)  
java.base/sun.nio.ch.NioSocketImpl.accept(NioSocketImpl.java:752)  
java.base/java.net.ServerSocket.implAccept(ServerSocket.java:649)  
java.base/java.net.ServerSocket.platformImplAccept(ServerSocket.java:615)  
java.base/java.net.ServerSocket.implAccept(ServerSocket.java:591)  
java.base/java.net.ServerSocket.implAccept(ServerSocket.java:548)  
java.base/java.net.ServerSocket.accept(ServerSocket.java:505)  
Main.lambda$main$0(Main.java:80)  
java.base/java.lang.Fiber.lambda$new$0(Fiber.java:195)  
java.base/java.lang.Continuation.enter(Continuation.java:372)  
java.base/java.lang.Continuation.run(Continuation.java:328)  
java.base/java.lang.Fiber.runContinuation(Fiber.java:366)  
java.base/java.util.concurrent.ForkJoinTask$RunnableExecuteAction.exec(ForkJoinTask.java:1425)  
java.base/java.util.concurrent.ForkJoinTask.doExec(ForkJoinTask.java:290)  
java.base/java.util.concurrent.ForkJoinPool$WorkQueue topLevelExec(ForkJoinPool.java:1017)  
java.base/java.util.concurrent.ForkJoinPool.scan(ForkJoinPool.java:1666)  
java.base/java.util.concurrent.ForkJoinPool.runWorker(ForkJoinPool.java:1599)  
java.base/java.util.concurrent.ForkJoinWorkerThread.run(ForkJoinWorkerThread.java:189)
```

```
var listener = new java.net.ServerSocket(...);  
listener.accept();
```



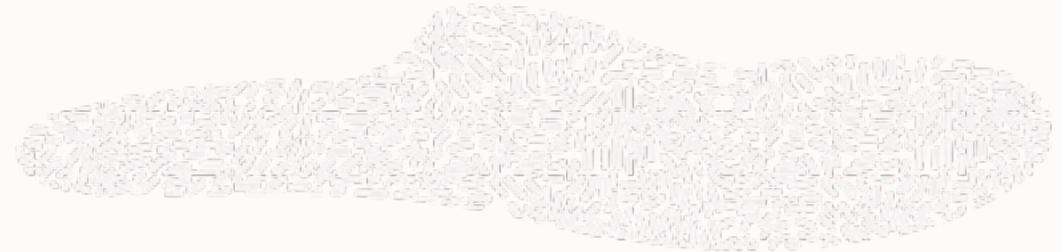
```
java.base/java.lang.Continuation.run(Continuation.java:328)  
java.base/java.lang.Fiber.runContinuation(Fiber.java:366)  
java.base/java.util.concurrent.ForkJoinTask$RunnableExecuteAction.exec(ForkJoinTask.java:1425)  
java.base/java.util.concurrent.ForkJoinTask.doExec(ForkJoinTask.java:290)  
java.base/java.util.concurrent.ForkJoinPool$WorkQueue topLevelExec(ForkJoinPool.java:1017)  
java.base/java.util.concurrent.ForkJoinPool.scan(ForkJoinPool.java:1666)  
java.base/java.util.concurrent.ForkJoinPool.runWorker(ForkJoinPool.java:1599)  
java.base/java.util.concurrent.ForkJoinWorkerThread.run(ForkJoinWorkerThread.java:189)
```

# Continuation: Limitations

- Can't yield with native frames on continuation stack (rare)
- Can't yield while holding a monitor (temporary)
- Virtual threads block their carrier when their continuation is pinned to preserve thread semantics.

```
class Continuation implements Runnable {  
    public Continuation(ContinuationScope scope, Runnable body);  
  
    public final void run();  
    public final boolean isDone();  
    public static void yield(ContinuationScope scope);  
  
    protected static Continuation currentContinuation(ContinuationScope  
scope);  
  
    public StackWalker stackWalker();  
  
    protected void onPinned(Reason reason)  
    { throw new IllegalStateException("Pinned: " + reason); }  
}
```

```
new Continuation (A, () -> {
    System.out.println("A begin");
    new Continuation(B, () -> {
        System.out.println("B begin");
        new Continuation(C, () -> {
            System.out.println("C begin");
            Continuation.yield(B);
            System.out.println("C end");
        }).run();
        System.out.println("B end");
    }).run();
    System.out.println("A end");
}).run();
```



A begin  
B begin  
C begin  
A end

# Implementation

- Continuation stacks are stored on the heap
- Not GC roots
- Copying to/from stack is *very cheap* (same no. of cache misses as an sp switch)
- Lazy-copying to/from stack with a return barrier (freeze/thaw)
- Finding references is *very expensive*

# Implementation

- Stack portions are copied into “stack chunk” objects without parsing.
- GC parses stack and finds oopmaps
- Chunks are mutable while they’re in a non-barrier region (young-gen and other special cases, depending on GC)
- Once a chunk object is promoted, frames can be thawed (removed) but not frozen (added). A new chunk is allocated on next freeze.
- Read-only chunks are trimmed.

Thank you

---



ORACLE

O